


Does a concentrated, ES-targeted Stage-C budget train better than a random one of the same size? Decided by a rule frozen before run 1: $\Delta > 2\hat{\sigma}$ and sign consistent 3/3 seeds.


SETUP



total_step pinned **253 / 127** in both rounds
of every run — added data buys **0** gradient steps

MEASUREMENT

SINet | COD10K-test (primary), S_{α} arm means:



A0 0.7010 · A2 0.7077 · B 0.7134 · C10 0.7185 — every arm lies *inside* the bar, and the whole A0→C10 span (0.0175) is narrower than the bar itself

the pre-registered claim

$\Delta(C-B) = +0.005034$, sign 3/3, against a bar of 0.017933

OUTCOME

$\Delta(C-B)$ is **WITHIN NOISE** in **all four** arch × endpoint cells, and the two architectures **disagree** on its sign (SINet +0.005034, SINet-v2 -0.000399).

The bar $2\hat{\sigma} = 0.017933$ is *larger* than the entire reference improvement of 0.0142 this design set out to half-resolve.

no effect resolvable at this sensitivity

a null bounded by power —
not a demonstration of absence