


Is the CHAMELEON evaluation set re-encoded copies of the data these models train on? And does the detector find contamination everywhere it looks, or can it return a clean answer?

(a) THE DETECTOR — exact hashing is the wrong instrument



DESC=32 RMS<=14 TOL=6.0

exact hashing calls CHAMELEON *perfectly clean* while 53.9% of it is re-encoded training data

(b) 41 IS A PROPERTY OF THE DATA, NOT OF THE CUTOFF

tolerance sweep saturates:



tolerance

a *second, independent* instrument agrees on the same integer: the sorted nearest-distance list has **41 below 5.51** and the next at **40.58** — a **7.36×** jump, and the author-sourced copy is byte-identical (76/76).

(c) THE NEGATIVE CONTROL — the same detector, imported unchanged



and the honesty bound: 3039 of 4121 NC4K images (73.7%) have no same-dimension candidate at all, so they are **unchecked, not clean**. Every rate here is a *lower bound*.

OUTCOME

CHAMELEON is **not an independent endpoint** for COD10K-trained models — and that is the whole claim.

D2R measured for a score-inflation figure and **found none** (all eight percentiles in 0.448–0.559, and the sign flips with the mask release), so none is reported.

independence, not inflation
a protocol violation does not need a score advantage to be disqualifying